

Introduction to **Atomic**: Tailoring a Trusted OS for Containers

Joe Brockmeier

jzb@redhat.com

@jzb

Introduction

- What is Project Atomic?
- Anatomy of an Atomic Host
- Coming Soon
- Getting Involved



**(I don't need to explain
containers, right? Good.)**

What is Project Atomic?



Project Atomic 101

- Upstream community for developing tools and patterns for developing Atomic hosts.
- Umbrella project for Red Hat's efforts around developing, building, running, and managing containers.
- *Not* a new distribution – Atomic Hosts are built from CentOS, Fedora, or Red Hat Enterprise Linux.



Why Atomic?

- We can run Linux containers on CentOS, Fedora, and RHEL already!
- Provide a streamlined host optimized for running and managing containers.
- All applications should be deployed as containers, rather than installing on the host.
- Host should be “cattle” and updates should be easy to deploy and manage.

What Atomic Hosts Provide

- Streamlined host based on CentOS, Fedora, or RHEL packages + container stack.
- rpm-ostree
- /usr/bin/atomic
- Docker
- Kubernetes
- Cockpit
- Super Privileged Containers (SPC)

What Atomic Hosts **Won't** Provide

- Atomic hosts are “immutable” – don't expect to install packages on running systems
- Official images are **minimal** – that means your favorite tool probably won't be added
 - *Aside from Atomic development or troubleshooting, you should never be logged into an Atomic Host*
- More than necessary

CentOS, Fedora, or RHEL?

- Aside from rpm-ostree, all of the components that make up an Atomic Host are shared w/the parent distribution.
- You want support? Go RHEL Atomic Host.
- CentOS Atomic is currently under development, and hasn't released any “official” images.
- Fedora 21 released in December – developed by the Cloud Working Group.
- A CentOS rebuild of RHEL AH is coming soon.



rpm-ostree's history

- OSTree initially developed for GNOME continuous by Colin Walters
- The rpm-ostree stuff came slightly later
- “Git for operating systems”
 - bootable, immutable, & versioned filesystem trees
 - works on top of any *nix filesystem
 - support for UID/GID, extended attr, handling bootloader, and more.

Why rpm-ostree?

- “Atomic” updates make more sense for an immutable system
- Preserves the tooling to create packages, allows re-use of RPMs rather than re-inventing the wheel
- Easy rollback in the event you need to return to known-good tree
- Clean transaction for updates

How rpm-ostree works (high level)

- Filesystem is read-only, except /var and /etc
- /etc is 3-way merged when you do an update
- All data (e.g. containers) is unchanged on upgrade
- Problem with an upgrade? `rpm-ostree rollback`

/usr/bin/atomic

- Coherent entry point to the system: manage host and containers with the atomic command.
- Fill gaps in Linux container implementations.
 - e.g. “atomic install foo” can install a container with its k8s configuration and/or systemd unit file.
 - “atomic run” grabs the LABEL “run” with its Docker cmd line. Saves the user much typing.
- The “atomic host” command can be used for rpm-ostree updates.

Cockpit

- Cockpit started prior to Atomic
- Server manager for administering Linux servers via the Web browser
- Doesn't interfere with normal admin tools
- Designed to be multi-server
- Support for managing containers, Kubernetes
- <http://cockpit-project.org/>



alpo

Dashboard

Cluster

System

Services

Containers >

Journal

Networking

Storage

Kubernetes Tool

Administrator Accounts

Terminal

Containers » compassionate_curie

Container: compassionate_curie

Start

Stop

Restart

Delete

Commit

Id: d72052b2d11a456cbd10fb230328e26f317fff9082a3bbdcbaa96581cd27cb8a


Created: 1428511989

Image: fedora:22

Command: /bin/bash

State: Up since 2015-04-17T07:16:52.481219236Z

Memory usage:



0.1 / 8.0 EB

CPU usage: 0%

1024 shares

[Change resource limits](#)

```
top - 07:17:08 up 7:52, 0 users, load average: 0.48, 0.23, 0.23
Tasks: 2 total, 1 running, 1 sleeping, 0 stopped, 0 zombie
%Cpu(s): 4.5 us, 0.8 sy, 0.0 ni, 94.0 id, 0.7 wa, 0.0 hi, 0.0 si, 0.0 st
KiB Mem : 7860492 total, 255048 free, 4163812 used, 3441632 buff/cache
KiB Swap: 7864316 total, 7858920 free, 5396 used, 1638984 avail Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
1	root	20	0	12540	3656	2576	S	0.0	0.0	0:00.03	bash
19	root	20	0	52372	3760	3268	R	0.0	0.0	0:00.00	top

Changes to 'docker search' & 'docker pull'

- We *mostly* ship vanilla Docker
- Additional registries for 'docker search' & 'docker pull'
- We add the RHEL registry to grab official RHEL content*
- Docker search lists fully qualified image name
- Ability to block registries
- Can warn on “push” to ensure private images aren't pushed to public registry

Super-Privileged Containers (SPC)

- We mean it when we say “run everything in containers” on Atomic
- Usually containers have limited interaction w/the host
- SPC containers can be run with ``atomic run`` which saves the need for long docker commands to enable privileges

Shipping Super-Privileged Containers (SPC)

- **RHEL Atomic Tools Container Image** – debugging tools like strace, traceroute, man pages, etc. needed to troubleshoot an image.
- **RHEL Atomic rsyslog Container Image** – runs rsyslogd service to send logs to central server, etc. (journald collects data either way)
- **RHEL Atomic sadc Container Image** – runs sadc from sysstat to be used w/ `sar`
- More to come!

Nulecule (in early development)

- Specification for multi-container application w/dependencies (“Atomic App”)
- Lets developer describe application, sysadmin define parameters for app at runtime
- Creates super-orchestration parameters for Kubernetes
- Defines on-demand scheduling of resource utilization
- Basis for policy-based orchestration via Mesos
- Supports Docker, ACI and potentially other container formats
- github.com/projectatomic/nulecule

Kubernetes

- Initially used GearD from OpenShift, phased out in favor of Kubernetes
- Working with upstream to improve / develop Kubernetes for container management

Pulling the Pieces Together



Fedora Atomic Hosts

- Work is being done through the Cloud Work Group & will be part of the Cloud Product
- First release in Fedora 21
- Adding new image formats in Fedora 22, updated Cockpit, etc.
- Moving to 2-week release cycle based on Rawhide or -current soon



CentOS 7 Atomic Hosts

- Work is being done through CentOS Atomic SIG
- CentOS-based Atomic Hosts are still in development, working out a few details like signing
- Will be providing a rebuild of RHEL Atomic Host soon
- CentOS SIG / Project Atomic will be providing a faster-moving release with packages in development soon



Getting Involved

- Website: projectatomic.io
- Github: github.com/projectatomic
- Facebook.com/[projectatomic](https://www.facebook.com/projectatomic)
- Twitter: [@projectatomic](https://twitter.com/projectatomic)
- Mailing Lists:
<http://www.projectatomic.io/community/>



Thank you!

jzb@redhat.com

Twitter: @jzb

@projectatomic

